

Problem Statement: Pick up the following stocks and generate forecasts accordingly Stocks:

1.NASDAQ.AAPL

2.NASDAQ.ADP

3.NASDAQ.CBOE

4.NASDAQ.CSCO

5.NASDAQ.EBAY

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from pandas.tools.plotting import autocorrelation_plot
from statsmodels.graphics.tsaplots import plot_pacf
from statsmodels.tsa.arima_model import ARIMA, ARMAResults
import datetime
import sys
import seaborn as sns
import statsmodels
import statsmodels.stats.diagnostic as diag
from statsmodels.tsa.stattools import adfuller
from scipy.stats.mstats import normaltest
from matplotlib.pyplot import acorr
plt.style.use('fivethirtyeight')
%matplotlib inline
```

```
In [2]: #Load data
df = pd.read_csv('C:/Users/HP/data_stocks.csv')
```

```
Out[2]:
```

	DATE	SP500	NASDAQ.AAL	NASDAQ.AAPL	NASDAQ.ADBE	NASDAQ.ADI	NASDAQ.ADP	NASDAQ..
0	1491226200	2363.6101	42.3300	143.6800	129.6300	82.040	102.2300	85
1	1491226260	2364.1001	42.3600	143.7000	130.3200	82.080	102.1400	85
2	1491226320	2362.6799	42.3100	143.6901	130.2250	82.030	102.2125	85
3	1491226380	2364.3101	42.3700	143.6400	130.0729	82.000	102.1400	85
4	1491226440	2364.8501	42.5378	143.6600	129.8800	82.035	102.0600	85

5 rows × 502 columns

```
In [3]: # display five records in the dataframe
```

```
Out[3]:
```

	0	1	2	3	4
DATE	1.491226e+09	1.491226e+09	1.491226e+09	1.491226e+09	1.491226e+09
SP500	2.363610e+03	2.364100e+03	2.362680e+03	2.364310e+03	2.364850e+03
NASDAQ.AAL	4.233000e+01	4.236000e+01	4.231000e+01	4.237000e+01	4.253780e+01
NASDAQ.AAPL	1.436800e+02	1.437000e+02	1.436901e+02	1.436400e+02	1.436600e+02
NASDAQ.ADBE	1.296300e+02	1.303200e+02	1.302250e+02	1.300729e+02	1.298800e+02
NASDAQ.ADI	8.204000e+01	8.208000e+01	8.203000e+01	8.200000e+01	8.203500e+01
NASDAQ.ADP	1.022300e+02	1.021400e+02	1.022125e+02	1.021400e+02	1.020600e+02
NASDAQ.ADSK	8.522000e+01	8.565000e+01	8.551000e+01	8.548720e+01	8.570010e+01
NASDAQ.AKAM	5.976000e+01	5.984000e+01	5.979500e+01	5.962000e+01	5.962000e+01
NASDAQ.ALXN	1.215200e+02	1.214800e+02	1.219300e+02	1.214400e+02	1.216000e+02
NASDAQ.AMAT	3.899000e+01	3.901000e+01	3.891000e+01	3.884000e+01	3.893000e+01
NASDAQ.AMD	1.461000e+01	1.471000e+01	1.464000e+01	1.463000e+01	1.467000e+01
NASDAQ.AMGN	1.646300e+02	1.646800e+02	1.649050e+02	1.647600e+02	1.648500e+02
NASDAQ.AMZN	8.885500e+02	8.871173e+02	8.875110e+02	8.862700e+02	8.865800e+02
NASDAQ.ATVI	4.985000e+01	4.994000e+01	4.986000e+01	4.991500e+01	4.991500e+01
NASDAQ.AVGO	2.191100e+02	2.199800e+02	2.193900e+02	2.193000e+02	2.191800e+02
NASDAQ.BBBY	3.943000e+01	3.968000e+01	3.960000e+01	3.957000e+01	3.955000e+01
NASDAQ.BIIB	2.740800e+02	2.739900e+02	2.742750e+02	2.735900e+02	2.735400e+02
NASDAQ.CA	3.178000e+01	3.178000e+01	3.176500e+01	3.183000e+01	3.183000e+01
NASDAQ.CBOE	8.103000e+01	8.121000e+01	8.121000e+01	8.113000e+01	8.112000e+01
NASDAQ.CELG	1.248900e+02	1.249900e+02	1.250000e+02	1.247300e+02	1.248300e+02
NASDAQ.CERN	5.882000e+01	5.849500e+01	5.847000e+01	5.842000e+01	5.860000e+01
NASDAQ.CHRW	7.772500e+01	7.794000e+01	7.781500e+01	7.795000e+01	7.805000e+01
NASDAQ.CHTR	3.307300e+02	3.307300e+02	3.307300e+02	3.307300e+02	3.307300e+02
NASDAQ.CINF	7.243000e+01	7.204000e+01	7.205500e+01	7.214000e+01	7.221500e+01
NASDAQ.CMCSA	3.747000e+01	3.754000e+01	3.761000e+01	3.762000e+01	3.762500e+01
NASDAQ.CME	1.193850e+02	1.188100e+02	1.188300e+02	1.186800e+02	1.189350e+02
NASDAQ.COST	1.677400e+02	1.677760e+02	1.680000e+02	1.682000e+02	1.680400e+02
NASDAQ.CSCO	3.374000e+01	3.388000e+01	3.390000e+01	3.384990e+01	3.384000e+01
NASDAQ.CSX	4.664500e+01	4.661000e+01	4.688500e+01	4.670000e+01	4.685620e+01
...
NYSE.USB	5.162000e+01	5.158000e+01	5.146000e+01	5.138000e+01	5.143470e+01
NYSE.UTX	1.123600e+02	1.123600e+02	1.121300e+02	1.120100e+02	1.122300e+02
NYSE.V	8.935000e+01	8.935000e+01	8.916000e+01	8.906000e+01	8.910000e+01
NYSE.VAR	9.113000e+01	9.121000e+01	9.108000e+01	9.101500e+01	9.100000e+01
NYSE.VFC	5.521000e+01	5.512000e+01	5.509000e+01	5.521000e+01	5.532000e+01
NYSE.VLO	6.659000e+01	6.635500e+01	6.624170e+01	6.617000e+01	6.618000e+01
NYSE.VMC	1.201300e+02	1.201300e+02	1.203368e+02	1.203100e+02	1.203600e+02
NYSE.VNO	1.003500e+02	1.000300e+02	1.003900e+02	1.003900e+02	1.001100e+02

```
In [4]: # make a list of columns
stock_features = ['NASDAQ.AAPL', 'NASDAQ.ADP', 'NASDAQ.CBOE', 'NASDAQ.CSCO', 'NASDAQ.EBAY']
col_list = ['DATE'] + stock_features
df1 = df[col_list]
```

Out[4]:

	DATE	NASDAQ.AAPL	NASDAQ.ADP	NASDAQ.CBOE	NASDAQ.CSCO	NASDAQ.EBAY
0	1491226200	143.6800	102.2300	81.03	33.7400	33.3975
1	1491226260	143.7000	102.1400	81.21	33.8800	33.3950
2	1491226320	143.6901	102.2125	81.21	33.9000	33.4100
3	1491226380	143.6400	102.1400	81.13	33.8499	33.3350
4	1491226440	143.6600	102.0600	81.12	33.8400	33.4000

In [5]:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41266 entries, 0 to 41265
Data columns (total 6 columns):
DATE                41266 non-null int64
NASDAQ.AAPL        41266 non-null float64
NASDAQ.ADP         41266 non-null float64
NASDAQ.CBOE        41266 non-null float64
NASDAQ.CSCO        41266 non-null float64
NASDAQ.EBAY        41266 non-null float64
dtypes: float64(5), int64(1)
memory usage: 1.9 MB
```

```
In [6]: #Checking for null values if any
```

```
Out[6]: DATE                0
NASDAQ.AAPL              0
NASDAQ.ADP               0
NASDAQ.CBOE              0
NASDAQ.CSCO              0
NASDAQ.EBAY              0
dtype: int64
```

```
In [7]: df1 = df1.copy()
df1['DATE'] = pd.to_datetime(df1['DATE'])
```

Out[7]:

	DATE	NASDAQ.AAPL	NASDAQ.ADP	NASDAQ.CBOE	NASDAQ.CSCO	NASDAQ.EBAY
41261	1970-01-01 00:00:01.504209360	164.11	106.565	100.89	32.185	36.135
41262	1970-01-01 00:00:01.504209420	164.12	106.590	100.88	32.200	36.130
41263	1970-01-01 00:00:01.504209480	164.01	106.520	100.86	32.200	36.130
41264	1970-01-01 00:00:01.504209540	163.88	106.400	100.83	32.195	36.120
41265	1970-01-01 00:00:01.504209600	163.98	106.470	100.89	32.225	36.130

In [8]:

Out[8]:

	DATE	NASDAQ.AAPL	NASDAQ.ADP	NASDAQ.CBOE	NASDAQ.CSCO	NASDAQ.EBAY
0	1970-01-01 00:00:01.491226200	143.6800	102.2300	81.03	33.7400	33.3975
1	1970-01-01 00:00:01.491226260	143.7000	102.1400	81.21	33.8800	33.3950
2	1970-01-01 00:00:01.491226320	143.6901	102.2125	81.21	33.9000	33.4100
3	1970-01-01 00:00:01.491226380	143.6400	102.1400	81.13	33.8499	33.3350
4	1970-01-01 00:00:01.491226440	143.6600	102.0600	81.12	33.8400	33.4000

In [9]: `df1 = df1.copy()`

In [10]:

Out[10]:

	DATE	NASDAQ.AAPL	NASDAQ.ADP	NASDAQ.CBOE	NASDAQ.CSCO	NASDAQ.EBAY	Month
0	1970-01-01 00:00:01.491226200	143.6800	102.2300	81.03	33.7400	33.3975	1970-01-01
1	1970-01-01 00:00:01.491226260	143.7000	102.1400	81.21	33.8800	33.3950	1970-01-01
2	1970-01-01 00:00:01.491226320	143.6901	102.2125	81.21	33.9000	33.4100	1970-01-01
3	1970-01-01 00:00:01.491226380	143.6400	102.1400	81.13	33.8499	33.3350	1970-01-01
4	1970-01-01 00:00:01.491226440	143.6600	102.0600	81.12	33.8400	33.4000	1970-01-01

In [11]: `col_list = ['Month'] + stock_features
df2 = df1[col_list]`

Out[11]:

	Month	NASDAQ.AAPL	NASDAQ.ADP	NASDAQ.CBOE	NASDAQ.CSCO	NASDAQ.EBAY
0	1970-01-01	143.6800	102.2300	81.03	33.7400	33.3975
1	1970-01-01	143.7000	102.1400	81.21	33.8800	33.3950
2	1970-01-01	143.6901	102.2125	81.21	33.9000	33.4100
3	1970-01-01	143.6400	102.1400	81.13	33.8499	33.3350
4	1970-01-01	143.6600	102.0600	81.12	33.8400	33.4000

In [12]:

Out[12]:

```
Month          0  
NASDAQ.AAPL    0  
NASDAQ.ADP     0  
NASDAQ.CBOE    0  
NASDAQ.CSCO    0  
NASDAQ.EBAY    0  
dtype: int64
```

In [13]:

Out[13]:

	count	mean	std	min	25%	50%	75%	max
NASDAQ.AAPL	41266.0	150.453566	6.236826	140.160	144.640	149.9450	155.065	164.51
NASDAQ.ADP	41266.0	103.480398	4.424244	95.870	101.300	102.4400	104.660	121.77
NASDAQ.CBOE	41266.0	89.325485	5.746178	80.000	84.140	89.3150	93.850	101.35
NASDAQ.CSCO	41266.0	32.139336	0.985571	30.365	31.455	31.7733	32.790	34.49
NASDAQ.EBAY	41266.0	34.794506	1.099296	31.890	34.065	34.7700	35.610	37.46

In [14]: `final = df2.copy()`

In [15]: *#Time Series Forecasting for NASDAQ.AAPL*

```
df_AAPL = final[['Month',stock_features[0]]]
```

Out[15]:

	Month	NASDAQ.AAPL
0	1970-01-01	143.6800
1	1970-01-01	143.7000
2	1970-01-01	143.6901
3	1970-01-01	143.6400
4	1970-01-01	143.6600

In [16]: `df_AAPL.set_index('Month',inplace=True)`

Out[16]:

NASDAQ.AAPL	
Month	
1970-01-01	143.6800
1970-01-01	143.7000
1970-01-01	143.6901
1970-01-01	143.6400
1970-01-01	143.6600

In [17]:

Out[17]: DatetimeIndex(['1970-01-01', '1970-01-01', '1970-01-01', '1970-01-01',
 '1970-01-01', '1970-01-01', '1970-01-01', '1970-01-01',
 '1970-01-01', '1970-01-01',
 ...,
 '1970-01-01', '1970-01-01', '1970-01-01', '1970-01-01',
 '1970-01-01', '1970-01-01', '1970-01-01', '1970-01-01',
 '1970-01-01', '1970-01-01'],
 dtype='datetime64[ns]', name='Month', length=41266, freq=None)

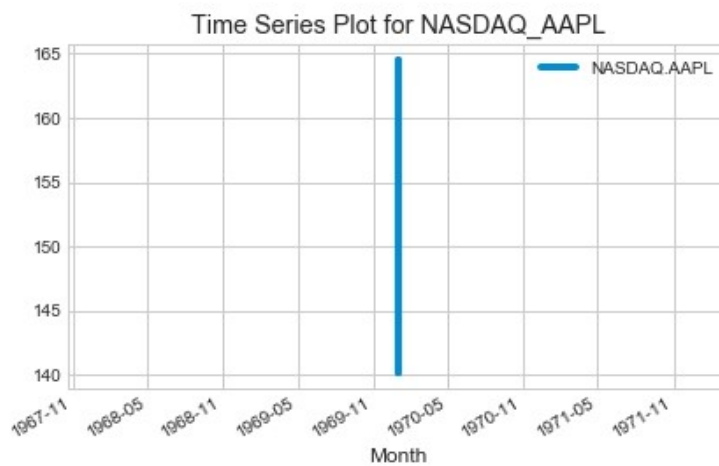
In [18]: *#Summary Statistics*

Out[18]:

	count	mean	std	min	25%	50%	75%	max
NASDAQ.AAPL	41266.0	150.453566	6.236826	140.16	144.64	149.945	155.065	164.51

```
In [19]: #Step 2 : Visualize the Data
```

```
import seaborn as sns
sns.set_style('whitegrid')
df_AAPL.plot()
plt.title('Time Series Plot for NASDAQ_AAPL')
```



```
In [20]: #Plotting Rolling Statistics and check for stationarity :
#The function will plot the moving mean or moving Standard Deviation. This is still vi

#NOTE: Moving mean and Moving standard deviation- At any instant 't',
```


```
In [21]: from statsmodels.tsa.stattools import adfuller
def test_stationarity(timeseries):

    #Determing rolling statistics
    rolmean = timeseries.rolling(12).mean()
    rolstd = timeseries.rolling(12).std()
    #Plot rolling statistics:
    plt.plot(timeseries, color='blue',label='Original')
    plt.plot(rolmean, color='red', label='Rolling Mean')
    plt.plot(rolstd, color='black', label = 'Rolling Std')
    plt.legend(loc='best')
    plt.title('Rolling Mean & Standard Deviation')
    plt.show()
    """
    Pass in a time series, returns ADF report
    """
    result = adfuller(timeseries)
    print('\nAugmented Dickey-Fuller Test:')
    labels = ['ADF Test Statistic', 'p-value', '#Lags Used', 'Number of Observations Used']

    for value,label in zip(result,labels):
        print(label+' : '+str(value) )
    for k,v in result[4].items():
        print('Critical {}: value {}'.format(k,v))

    if result[1] <= 0.05:
        print("strong evidence against the null hypothesis, reject the null hypothesis")
    else:
        print("weak evidence against null hypothesis, time series has a unit root, indicating it is non-stationary")

test_stationarity(df AAPL['NASDAQ.AAPL'])
```



```
Augmented Dickey-Fuller Test:
ADF Test Statistic : -0.9128532997926669
p-value : 0.7837101772613867
#Lags Used : 31
Number of Observations Used : 41234
Critical 1% : value -3.4305085998723857
Critical 5% : value -2.8616100975579815
Critical 10% : value -2.5668073106689477
weak evidence against null hypothesis, time series has a unit root, indicating it
is non-stationary
```

```
In [22]: #Note:This is not stationary because :Mean is increasing even though the Std is small.
#Test stat is > critical value.
```

```
In [23]: #MAKE THE TIME SERIES STATIONARY
#There are two factors that make a time series non-stationary. They are:

#Trend: non-constant mean
#Seasonality: Variation at specific time-frames

#Differencing
#The first difference of a time series is the series of changes from one period to the
#We can do this easily with pandas.
#You can continue to take the second difference, third difference, and so on until you
```

```
In [24]: #First Difference

df_AAPL = df_AAPL.copy()
df_AAPL.loc[:, 'First_Difference'] = df_AAPL['NASDAQ.AAPL'] - df_AAPL['NASDAQ.AAPL'].shift(1)
```


Out[24]:

NASDAQ.AAPL First_Difference		
Month		
1970-01-01	143.6800	NaN
1970-01-01	143.7000	0.0200
1970-01-01	143.6901	-0.0099
1970-01-01	143.6400	-0.0501
1970-01-01	143.6600	0.0200

```
In [25]: df_AAPL = df_AAPL.copy()
```



```
In [26]: #Test Staionarity
```



Augmented Dickey-Fuller Test:

ADF Test Statistic : -35.73774148340111

p-value : 0.0

#Lags Used : 30

Number of Observations Used : 41234

Critical 1% : value -3.4305085998723857

Critical 5% : value -2.8616100975579815


Critical 10% : value -2.5668073106689477

strong evidence against the null hypothesis, reject the null hypothesis. Data has no unit root and is stationary

```
In [27]: #Seasonal Decomposition


from statsmodels.tsa.seasonal import seasonal_decompose
plt.figure(figsize=(11,8))
decomposition = seasonal_decompose(df AAPL['NASDAQ.AAPL'],freq=12)
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid
plt.subplot(411)
plt.plot(df AAPL['NASDAQ.AAPL'],label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend,label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal,label='Seasonal')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual,label='Residual')
```

Out[27]: <matplotlib.legend.Legend at 0x1b689ccd518>



```
In [28]: #Note:
```

```
In [29]: ts_log_decompose = residual
ts_log_decompose.dropna(inplace=True)
```



```
Augmented Dickey-Fuller Test:
ADF Test Statistic : -43.043433535543166
p-value : 0.0
#Lags Used : 55
Number of Observations Used : 41197
Critical 1% : value -3.4305087423235587
Critical 5% : value -2.861610160516496
Critical 10% : value -2.566807344180027
strong evidence against the null hypothesis, reject the null hypothesis. Data has
no unit root and is stationary
```

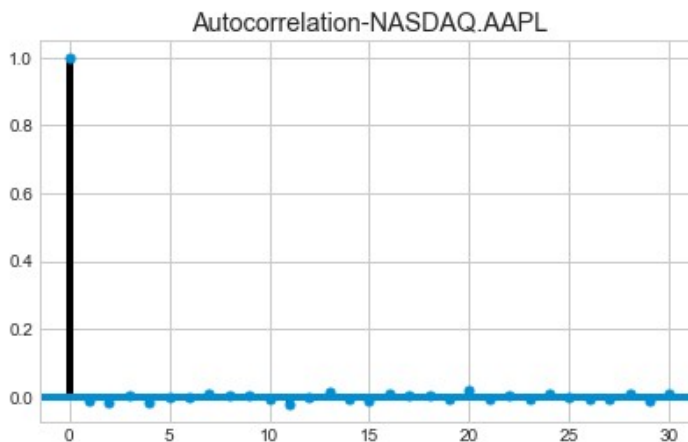
```
In [30]: #Note :This is stationary because:

#Test statistic is lower than critical values.
```

```
In [31]: #Autocorrelation and Partial Autocorrelation Plots

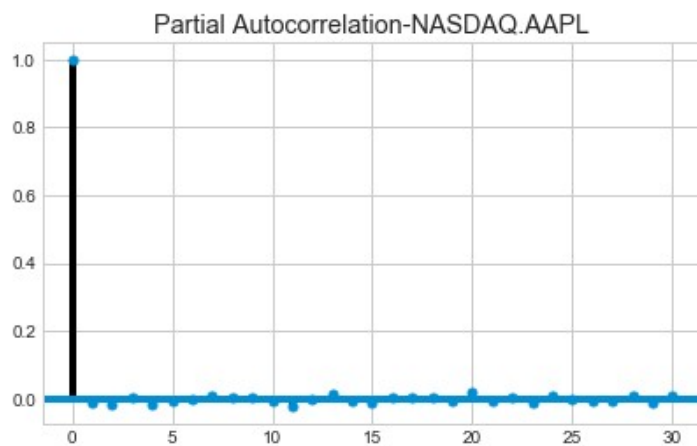
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
plt.figure(figsize=(20,8))

<Figure size 1440x576 with 0 Axes>
```




```
In [32]: plt.figure(figsize=(20,8))
```

<Figure size 1440x576 with 0 Axes>



```
In [33]: from pandas.plotting import autocorrelation_plot
```

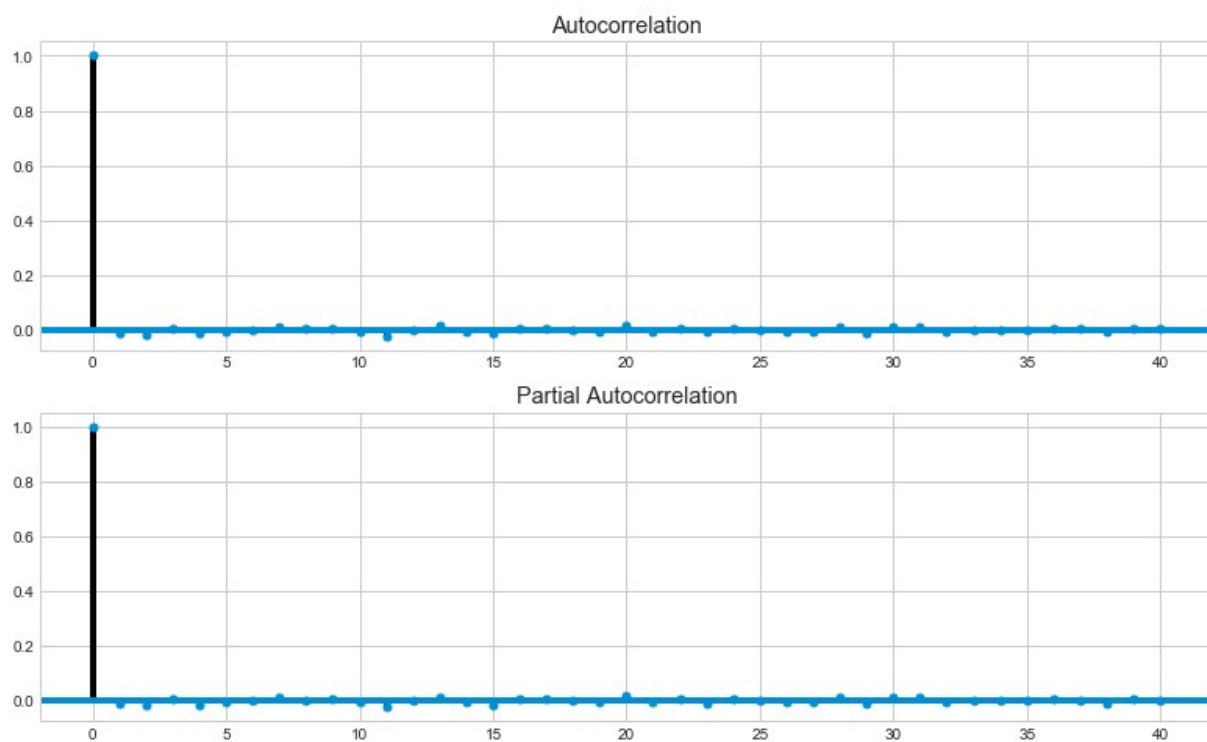
```
Out[33]: <matplotlib.axes._subplots.AxesSubplot at 0x1b689d1bfd0>
```



```
In [34]: #Forecasting a Time Series
#Auto Regressive Integrated Moving Average (ARIMA) -

import statsmodels.api as sm
from statsmodels.tsa.arima_model import ARIMA, ARIMAResults
from statsmodels.tsa.stattools import acf, pacf
```

```
In [35]: fig = plt.figure(figsize=(12,8))  
ax1 = fig.add_subplot(211)  
fig = sm.graphics.tsa.plot_acf(df AAPL['First_Difference'].iloc[30:], lags=40, ax=ax1)  
ax2 = fig.add_subplot(212)
```



```
In [36]: lag_acf = acf(df AAPL['First_Difference'], nlags=80)
```


```
In [37]: plt.figure(figsize=(10,10))
plt.subplot(121)
plt.plot(lag_acf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_AAPL['First_Difference'])),linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_AAPL['First_Difference'])),linestyle='--',color='gray')

plt.title('Autocorrelation')

plt.subplot(122)

plt.plot(lag_pacf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_AAPL['First_Difference'])),linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_AAPL['First_Difference'])),linestyle='--',color='gray')

plt.title('Partial Autocorrelation')
```



In []:

```
In [38]: #Note- The two dotted lines on either sides of 0 are the confidence intervals.

#These can be used to determine the 'p' and 'q' values as:

#p: The first time where the PACF crosses the upper confidence interval, here its close
```

```
In [39]: #Using the Seasonal ARIMA model

model= sm.tsa.statespace.SARIMAX(df_AAPL['NASDAQ.AAPL'],order=(0,1,0),seasonal_order=(
results = model.fit()
```

C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:225: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
' ignored when e.g. forecasting.', ValueWarning)

Statespace Model Results

```
=====
=====
Dep. Variable:          NASDAQ.AAPL    No. Observations:
41265
Model:                SARIMAX(0, 1, 0)x(0, 1, 0, 12)    Log Likelihood          2
4925.552
Date:                  Mon, 25 Feb 2019    AIC                  -4
9849.104
Time:                  14:56:08    BIC                  -4
9840.476
Sample:                0    HQIC                  -4
9846.377

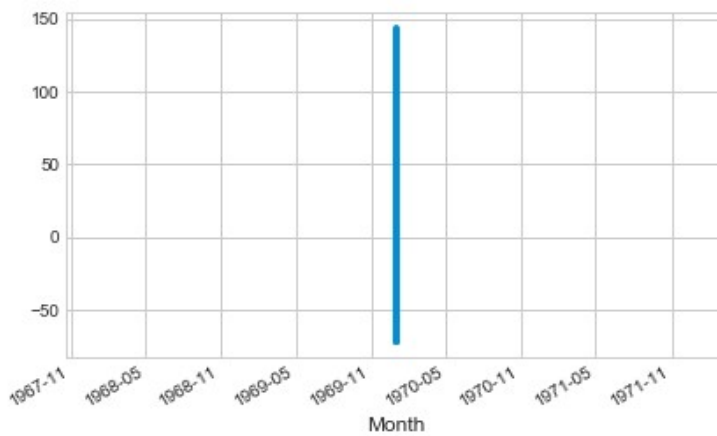
- 41265
Covariance Type:      opg
=====
=====
              coef      std err      z      P>|z|      [0.025      0.975]
-----
sigma2          0.0175    4.57e-06   3828.755    0.000      0.017      0.017
=====
=
Ljung-Box (Q):          10611.64    Jarque-Bera (JB):      3462262324.8
9
Prob(Q):                0.00    Prob(JB):                0.0
0
Heteroskedasticity (H): 2.92    Skew:                  -2.0
0
Prob(H) (two-sided):    0.00    Kurtosis:              1422.2
6
=====
=
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


In [40]:

Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0x1b68acaf710>



In [41]:

Out[41]: <matplotlib.axes._subplots.AxesSubplot at 0x1b68bd9beb8>




```
In [42]: df_AAPL = df_AAPL.copy()
df_AAPL['Forecast'] = results.predict()
```

Out[42]:

	NASDAQ.AAPL	First_Difference	Forecast
Month			
1970-01-01	143.7000	0.0200	0.0000
1970-01-01	143.6901	-0.0099	143.7000
1970-01-01	143.6400	-0.0501	143.6901
1970-01-01	143.6600	0.0200	143.6400
1970-01-01	143.7800	0.1200	143.6600


```
In [43]: #Prediction of Future Values
```

```
Out[43]: <matplotlib.axes._subplots.AxesSubplot at 0x1b68bdf7b00>
```



```
In [44]:
```

```
C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.
ValueWarning)
```

```
Out[44]: 41265    163.960
         41266    163.935
         41267    163.910
         41268    163.810
         41269    163.940
         41270    163.950
         41271    163.890
         41272    163.860
         41273    163.870
         41274    163.760
         dtype: float64
```

In [45]:

```
C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.
  ValueWarning)
```

```
Out[45]: 41264    163.930
         41265    163.960
         41266    163.935
         41267    163.910
         41268    163.810
         41269    163.940
         41270    163.950
         41271    163.890
         41272    163.860
         41273    163.870
         41274    163.760
         dtype: float64
```

In [46]:

```
C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.
  ValueWarning)
```

```
Out[46]: 41264    163.930
         41265    163.960
         41266    163.935
         41267    163.910
         41268    163.810
         41269    163.940
         41270    163.950
         41271    163.890
         41272    163.860
         41273    163.870
         41274    163.760
         dtype: float64
```

In [47]:

```
#Accuracy of the Forecast using MSE-Mean Squared Error

from sklearn.metrics import mean_squared_error, mean_absolute_error
print('Mean Squared Error NASDAQ.AAPL -', mean_squared_error(df_AAPL['NASDAQ.AAPL'], df_AAPL['NASDAQ.AAPL']))

Mean Squared Error NASDAQ.AAPL - 0.6426408212261927
Mean Absolute Error NASDAQ.AAPL - 0.07550728219978
```

In [48]:

```
#Time Series Forecasting for NASDAQ.ADP

df_ADP = final[['Month', stock_features[1]]]
```

Out[48]:

	Month	NASDAQ.ADP
0	1970-01-01	102.2300
1	1970-01-01	102.1400
2	1970-01-01	102.2125
3	1970-01-01	102.1400
4	1970-01-01	102.0600


```
In [49]: df_ADP.set_index('Month', inplace=True)
```

```
Out[49]:
```


NASDAQ.ADP	
Month	
1970-01-01	102.2300
1970-01-01	102.1400
1970-01-01	102.2125
1970-01-01	102.1400
1970-01-01	102.0600

```
In [50]: #Visualize Data

df_ADP.plot()
plt.title('Time Series Plot for NASDAQ_ADP')
```



In [51]:



Augmented Dickey-Fuller Test:

ADF Test Statistic : -1.7041735251574957

p-value : 0.42896344420667615

#Lags Used : 39

Number of Observations Used : 41226

Critical 1% : value -3.4305086306509716

Critical 5% : value -2.861610111161057

Critical 10% : value -2.5668073179094897

weak evidence against null hypothesis, time series has a unit root, indicating it is non-stationary

In [52]: *#MAKING THE TIME SERIES STATIONARY**#Differencing*

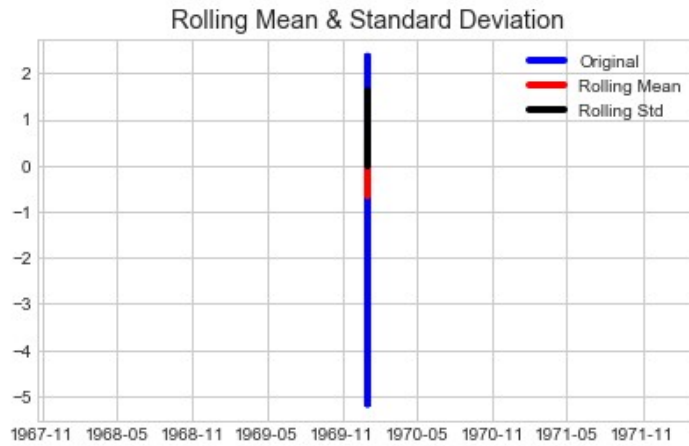
df_ADp = df_ADp.copy()

df_ADp['First_Difference'] = df_ADp['NASDAQ.ADP'] - df_ADp['NASDAQ.ADP'].shift(1)

Out[52]:

NASDAQ.ADP First_Difference		
Month		
1970-01-01	102.2300	NaN
1970-01-01	102.1400	-0.0900
1970-01-01	102.2125	0.0725
1970-01-01	102.1400	-0.0725
1970-01-01	102.0600	-0.0800

```
In [53]: df_ADp.dropna(inplace=True)
test_stationarity(df_ADp['First_Difference'])
```



Augmented Dickey-Fuller Test:

ADF Test Statistic : -31.055662244632316

p-value : 0.0

#Lags Used : 38

Number of Observations Used : 41226

Critical 1% : value -3.4305086306509716

Critical 5% : value -2.861610111161057


Critical 10% : value -2.5668073179094897

strong evidence against the null hypothesis, reject the null hypothesis. Data has no unit root and is stationary

In [54]: *#Seasonal Decomposition*


```
from statsmodels.tsa.seasonal import seasonal_decompose
plt.figure(figsize=(11,8))
decomposition = seasonal_decompose(df_ADP['First_Difference'],freq=12)
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid
plt.subplot(411)
plt.plot(df_ADP['First_Difference'],label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend,label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal,label='Seasonal')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual,label='Residual')
```

Out[54]: <matplotlib.legend.Legend at 0x1b68ac49400>



In [55]: *#Note: The data for NASDAQ.ADP is seasonal as interpreted from the seasonal plot of se*

```
In [56]: ts_log_decompose = residual
ts_log_decompose.dropna(inplace=True)
```




Augmented Dickey-Fuller Test:
 ADF Test Statistic : -57.848665441142764
 p-value : 0.0
 #Lags Used : 55
 Number of Observations Used : 41197
 Critical 1% : value -3.4305087423235587
 Critical 5% : value -2.861610160516496
 Critical 10% : value -2.566807344180027
 strong evidence against the null hypothesis, reject the null hypothesis. Data has no unit root and is stationary

```
In [57]: #Note : This is stationary because:

#Test statistic is lower than 1% critical values.
```

```
In [58]: #Autocorrelation and Partial Corelation plot

fig = plt.figure(figsize=(20,8))
ax1 = fig.add_subplot(211)
fig = sm.graphics.tsa.plot_acf(df_ADP['First_Difference'].iloc[38:], lags=80, ax=ax1)
ax2 = fig.add_subplot(212)
```



```
In [59]: lag_acf = acf(df_ADP['First_Difference'], nlags=80)
```

```

In [60]: plt.figure(figsize=(20,8))
plt.subplot(121)
plt.plot(lag_acf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_ADP['First_Difference'])),linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_ADP['First_Difference'])),linestyle='--',color='gray')

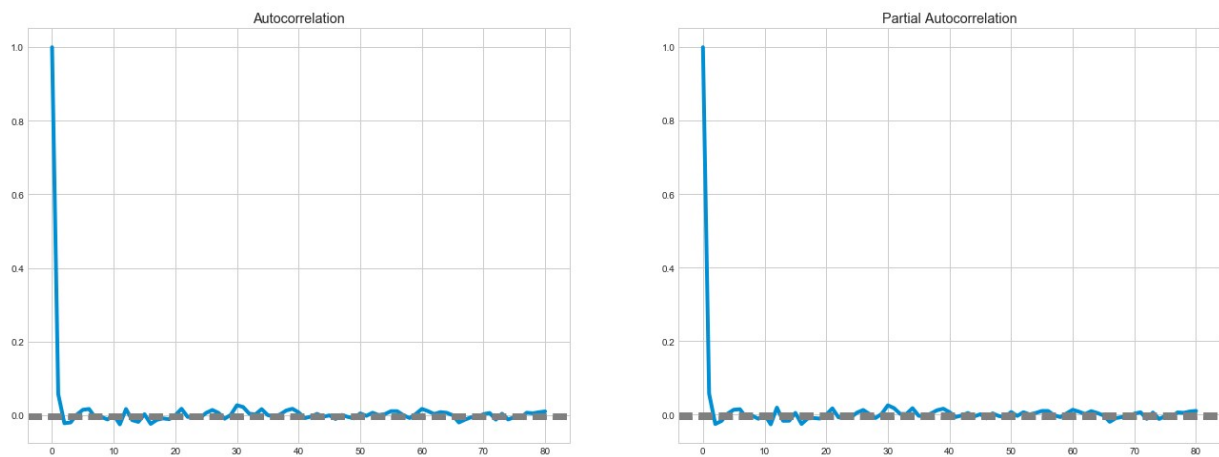
plt.title('Autocorrelation')

plt.subplot(122)

plt.plot(lag_pacf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_ADP['First_Difference'])),linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_ADP['First_Difference'])),linestyle='--',color='gray')

```

Out[60]: Text(0.5,1,'Partial Autocorrelation')



```

In [61]: #Note- The two dotted lines on either sides of 0 are the confidence intervals.

#These can be used to determine the 'p' and 'q' values as:

#p: The first time where the PACF crosses the upper confidence interval, here its close

```



```
In [62]: model = sm.tsa.statespace.SARIMAX(df_ADP['NASDAQ.ADP'], order=(0, 1, 0), seasonal_order=(0,
results = model.fit())
```

C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:225: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

' ignored when e.g. forecasting.', ValueWarning)

Statespace Model Results

```
=====
Dep. Variable:          NASDAQ.ADP    No. Observations:
41265
Model:          SARIMAX(0, 1, 0)x(0, 1, 0, 12)    Log Likelihood          3
4733.013
Date:           Mon, 25 Feb 2019    AIC          -6
9464.026
Time:           14:59:18    BIC          -6
9455.399
Sample:         0    HQIC          -6
9461.299
                    - 41265
Covariance Type:          opg
=====
              coef      std err      z      P>|z|      [0.025      0.975]
-----
sigma2          0.0109   5.34e-06  2036.708    0.000    0.011    0.011
=====
=
Ljung-Box (Q):          10628.96    Jarque-Bera (JB):          275266212.1
0
Prob(Q):          0.00    Prob(JB):          0.0
0
Heteroskedasticity (H):          2.20    Skew:          -1.5
9
Prob(H) (two-sided):          0.00    Kurtosis:          403.1
7
=====
=
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


```
In [63]:
```

```
Out[63]: [matplotlib.lines.Line2D at 0x1b6a12c3748]
```



```
In [64]: sns.set_style('whitegrid')
```

```
Out[64]: <matplotlib.axes._subplots.AxesSubplot at 0x1b6a1338978>
```



```
In [65]: df_ADP['Forecast'] = results.predict()
```

```
Out[65]:
```

	NASDAQ.ADP	Forecast
Month		
1970-01-01	106.565	106.705
1970-01-01	106.590	106.525
1970-01-01	106.520	106.510
1970-01-01	106.400	106.480
1970-01-01	106.470	106.430

```
In [66]:
```

```
C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.
  ValueWarning)
```

```
Out[66]: 41265    106.470
41266    106.470
41267    106.440
41268    106.380
41269    106.440
41270    106.420
41271    106.450
41272    106.385
41273    106.410
41274    106.340
dtype: float64
```


In [67]:

```
C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.
ValueWarning)
```

```
Out[67]: 41264      106.430
         41265      106.470
         41266      106.470
         41267      106.440
         41268      106.380
         41269      106.440
         41270      106.420
         41271      106.450
         41272      106.385
         41273      106.410
         41274      106.340
         41275      106.220
         dtype: float64
```

In [68]:

```
Out[68]: <matplotlib.axes._subplots.AxesSubplot at 0x1b6a132f550>
```



```
In [69]: from sklearn.metrics import mean_squared_error, mean_absolute_error
print('Mean Squared Error NASDAQ.AAPL -', mean_squared_error(df_ADP['NASDAQ.ADP'], df_Forecast['NASDAQ.ADP']))
Mean Squared Error NASDAQ.AAPL - 0.32679381130345103
Mean Absolute Error NASDAQ.AAPL - 0.0533967381818573
```



```
In [70]: #Times Series Forecasting for 'NASDAQ.CBOE'

df_CBOE= final[['Month',stock_features[2]]]
print(df_CBOE.head())
df_CBOE.set_index('Month',inplace=True)
print(df_CBOE.head())

df_CBOE.plot()
plt.title('Time Series Plot for NASDAQ_CBOE')
plt.show()
#test Stationarity
```

	Month	NASDAQ.CBOE
0	1970-01-01	81.03
1	1970-01-01	81.21
2	1970-01-01	81.21
3	1970-01-01	81.13
4	1970-01-01	81.12

	Month	NASDAQ.CBOE
	1970-01-01	81.03
	1970-01-01	81.21
	1970-01-01	81.21
	1970-01-01	81.13
	1970-01-01	81.12



```
Augmented Dickey-Fuller Test:
ADF Test Statistic : 0.16633930282615195
p-value : 0.9703092030510077
#Lags Used : 27
Number of Observations Used : 41238
```

```
In [71]: #MAKING THE TIME SERIES STATIONARY

#Differencing

df_CBOE = df_CBOE.copy()
```

Out[71]:

NASDAQ.CBOE	
Month	
1970-01-01	81.03
1970-01-01	81.21
1970-01-01	81.21
1970-01-01	81.13
1970-01-01	81.12


```
In [72]: df_CBOE['First_Difference'] = df_CBOE['NASDAQ.CBOE'] - df_CBOE['NASDAQ.CBOE'].shift(1)
```

Out[72]:

NASDAQ.CBOE First_Difference		
Month		
1970-01-01	81.03	NaN
1970-01-01	81.21	0.18
1970-01-01	81.21	0.00
1970-01-01	81.13	-0.08
1970-01-01	81.12	-0.01

In [73]:

```
In [74]: #Test Seasonality
```



Augmented Dickey-Fuller Test:

ADF Test Statistic : -41.642093645431416

p-value : 0.0

#Lags Used : 26

Number of Observations Used : 41238

Critical 1% : value -3.430508584487571

Critical 5% : value -2.8616100907584228

Critical 10% : value -2.5668073070497304

strong evidence against the null hypothesis, reject the null hypothesis. Data has no unit root and is stationary

```


In [75]: #Seasonal Decomposition
from statsmodels.tsa.seasonal import seasonal_decompose
plt.figure(figsize=(11,8))
decomposition = seasonal_decompose(df_CBOE['NASDAQ.CBOE'],freq=12)
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid
plt.subplot(411)
plt.plot(df_CBOE['NASDAQ.CBOE'],label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend,label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal,label='Seasonal')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual,label='Residual')

```

Out[75]: <matplotlib.legend.Legend at 0x1b68ac526d8>



```
In [76]: ts_log_decompose = residual
ts_log_decompose.dropna(inplace=True)
```



Augmented Dickey-Fuller Test:

ADF Test Statistic : -46.21672053215934

p-value : 0.0

#Lags Used : 55

Number of Observations Used : 41197

Critical 1% : value -3.4305087423235587


Critical 5% : value -2.861610160516496

Critical 10% : value -2.566807344180027

strong evidence against the null hypothesis, reject the null hypothesis. Data has no unit root and is stationary

```
In [77]: #Autocorrelation and Partial Corelation plot
```

```
fig = plt.figure(figsize=(20,8))
ax1 = fig.add_subplot(211)
fig = sm.graphics.tsa.plot_acf(df_CBOE['First_Difference'].iloc[26:], lags=80, ax=ax1)
ax2 = fig.add_subplot(212)
```




```
In [78]: lag_acf = acf(df_CBOE['First_Difference'],nlags=80)
lag_pacf = pacf(df_CBOE['First_Difference'],nlags=80,method='ols')

plt.figure(figsize=(11,8))
plt.subplot(121)
plt.plot(lag_acf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_CBOE['First_Difference'])),linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_CBOE['First_Difference'])),linestyle='--',color='gray')

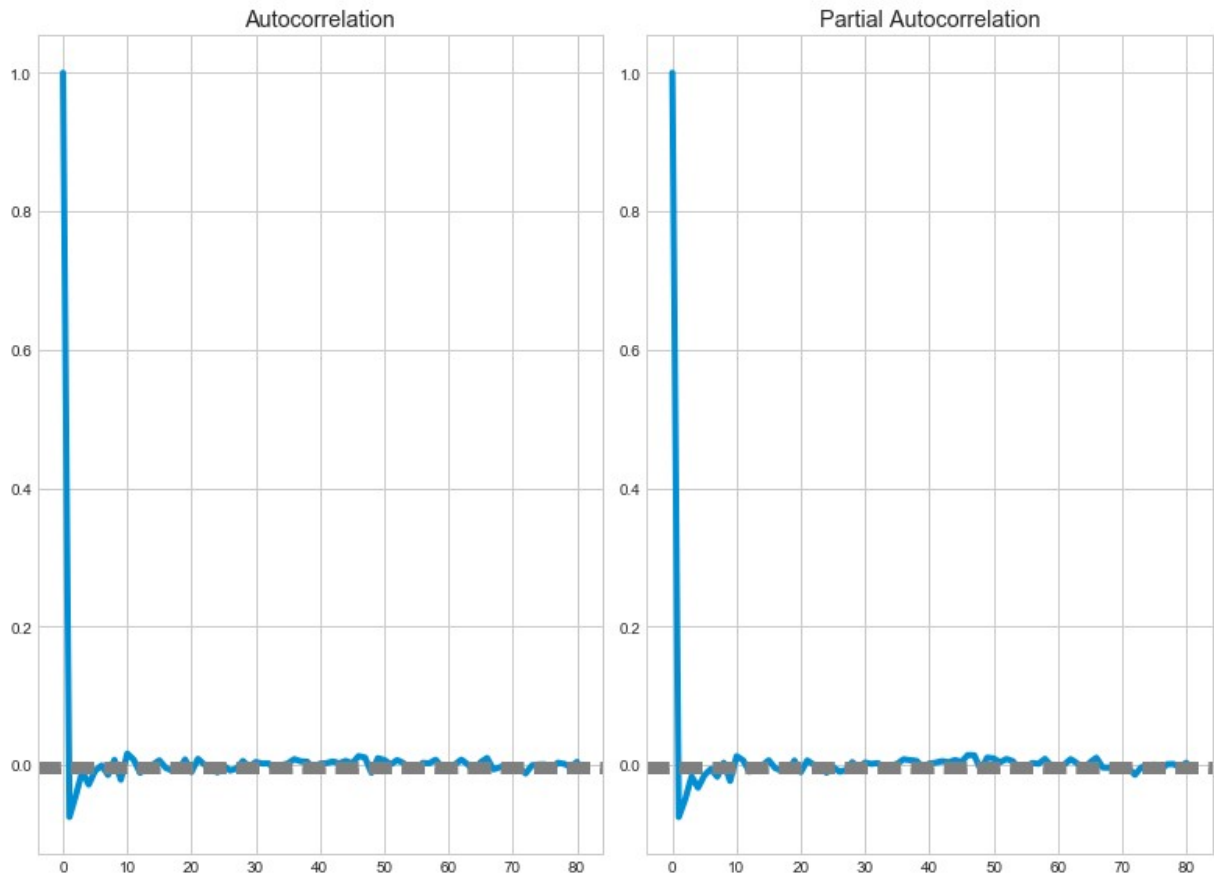
plt.title('Autocorrelation')

plt.subplot(122)

plt.plot(lag_pacf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_CBOE['First_Difference'])),linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_CBOE['First_Difference'])),linestyle='--',color='gray')

plt.title('Partial Autocorrelation')

plt.tight_layout()
```



```
In [79]: #Note- The two dotted lines on either sides of 0 are the confidence intervals.

#These can be used to determine the 'p' and 'q' values as:

#p: The first time where the PACF crosses the upper confidence interval, here its close
```

```
In [80]: # fit model
model= sm.tsa.statespace.SARIMAX(df_CBOE['NASDAQ.CBOE'],order=(0,1,0),seasonal_order=(
results = model.fit()
print(results.summary())
print(results.forecast())
df_CBOE['Forecast'] = results.predict()
df_CBOE[['NASDAQ.CBOE','Forecast']].plot(figsize=(20,8))
```

C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:225: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

' ignored when e.g. forecasting.', ValueWarning)

C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\base\model.py:508: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
"Check mle_retvals", ConvergenceWarning)

Statespace Model Results

```
=====
=====
Dep. Variable:                NASDAQ.CBOE    No. Observations:
41265
Model:                SARIMAX(0, 1, 0)x(0, 1, 0, 12)    Log Likelihood                5
3414.092
Date:                Mon, 25 Feb 2019    AIC                -10
6826.185
Time:                15:02:20    BIC                -10
6817.557
Sample:                0    HQIC                -10
6823.458

- 41265
Covariance Type:                opg
=====
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
sigma2          0.0044    5.33e-06    824.256      0.000      0.004      0.004
=====
=
Ljung-Box (Q):                11084.06    Jarque-Bera (JB):                7011759.6
5
Prob(Q):                0.00    Prob(JB):                0.0
0
Heteroskedasticity (H):                0.94    Skew:                -0.4
6
Prob(H) (two-sided):                0.00    Kurtosis:                66.8
6
=====
=
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.

ValueWarning)

41265 100.84

dtype: float64



In [81]:

```
C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.
  ValueWarning)
```

```
Out[81]: 41265    100.8400
         41266    100.8900
         41267    100.9100
         41268    100.8700
         41269    100.8800
         41270    100.8700
         41271    100.8799
         41272    100.8800
         41273    100.8700
         41274    100.8500
         dtype: float64
```

In [82]:

```
C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.
  ValueWarning)
```

```
Out[82]: 41264    100.8200
         41265    100.8400
         41266    100.8900
         41267    100.9100
         41268    100.8700
         41269    100.8800
         41270    100.8700
         41271    100.8799
         41272    100.8800
         41273    100.8700
         dtype: float64
```

```
In [83]: from sklearn.metrics import mean_squared_error, mean_absolute_error
         print('Mean Squared Error NASDAQ.CBOE -', mean_squared_error(df_CBOE['NASDAQ.CBOE'], df
```



```
Mean Squared Error NASDAQ.CBOE - 0.20399400190180045
Mean Absolute Error NASDAQ.CBOE - 0.04356630532804115
```

```
In [84]: #Time Series ForeCasting for 'NASDAQ.CSCO'

df_CSCO = final[['Month',stock_features[3]]]
print(df_CSCO.head())
df_CSCO.set_index('Month',inplace=True)
print(df_CSCO.head())
df_CSCO.plot()
plt.title("Time Series Plot for NASDAQ.CSCO")
plt.show()
#Test Staionarity
```

	Month	NASDAQ.CSCO
0	1970-01-01	33.7400
1	1970-01-01	33.8800
2	1970-01-01	33.9000
3	1970-01-01	33.8499
4	1970-01-01	33.8400

	Month	NASDAQ.CSCO
1970-01-01	33.7400	
1970-01-01	33.8800	
1970-01-01	33.9000	
1970-01-01	33.8499	
1970-01-01	33.8400	




Augmented Dickey-Fuller Test:
 ADF Test Statistic : -2.395554610889472
 p-value : 0.14299501995164166
 #Lags Used : 47
 Number of Observations Used : 41218
 Critical 1% : value -3.430508661441506

```
In [85]: #MAKING TIME SERIES STATIONARY

#Differencing

df_CSCO = df_CSCO.copy()
df_CSCO['First_Difference'] = df_CSCO['NASDAQ.CSCO'] - df_CSCO['NASDAQ.CSCO'].shift(1)
df_CSCO.dropna(inplace=True)
```




```
Augmented Dickey-Fuller Test:
ADF Test Statistic : -30.356682532566367
p-value : 0.0
#Lags Used : 46
Number of Observations Used : 41218
Critical 1% : value -3.430508661441506
Critical 5% : value -2.8616101247694137
Critical 10% : value -2.566807325152842
strong evidence against the null hypothesis, reject the null hypothesis. Data has
no unit root and is stationary
```

```


In [86]: #Seasonal Decomposition
from statsmodels.tsa.seasonal import seasonal_decompose
plt.figure(figsize=(11,8))
decomposition = seasonal_decompose(df_CSCO['NASDAQ.CSCO'],freq=12)
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid
plt.subplot(411)
plt.plot(df_CSCO['NASDAQ.CSCO'],label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend,label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal,label='Seasonal')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual,label='Residual')

```

Out[86]: <matplotlib.legend.Legend at 0x1b6ca5bcdd8>



```
In [87]: ts_log_decompose = residual
         ts_log_decompose.dropna(inplace=True)
```



Augmented Dickey-Fuller Test:
 ADF Test Statistic : -43.94517780543432
 p-value : 0.0
 #Lags Used : 55
 Number of Observations Used : 41197
 Critical 1% : value -3.4305087423235587
 Critical 5% : value -2.861610160516496
 Critical 10% : value -2.566807344180027
 strong evidence against the null hypothesis, reject the null hypothesis. Data has no unit root and is stationary


```
In [88]: #Note : This is stationary because:

         #Test statistic is lower than critical values.

         #The mean and std variations have small variations with time
```

```
In [89]: #Auto Corealtion and Partial Autocorelation Plots

         fig = plt.figure(figsize=(20,8))
         ax1 = fig.add_subplot(211)
         fig = sm.graphics.tsa.plot_acf(df_CSCO['First_Difference'].iloc[46:], lags=80, ax=ax1)
         ax2 = fig.add_subplot(212)
```



```
In [90]: lag_acf = acf(df_CSCO['First_Difference'],nlags=80)
```

```

In [91]: plt.figure(figsize=(20,8))
plt.subplot(121)
plt.plot(lag_acf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_CSC0['First_Difference'])),linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_CSC0['First_Difference'])),linestyle='--',color='gray')

plt.title('Autocorrelation')


plt.subplot(122)

plt.plot(lag_pacf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_CSC0['First_Difference'])),linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_CSC0['First_Difference'])),linestyle='--',color='gray')

plt.title('Partial Autocorrelation')

plt.tight_layout()

```



```

In [92]: #Note- The two dotted lines on either sides of 0 are the confidence intervals.

#These can be used to determine the 'p' and 'q' values as:

#p: The first time where the PACF crosses the upper confidence interval, here its close

```



```
In [93]: # fit model
model= sm.tsa.statespace.SARIMAX(df_CSCO['NASDAQ.CSCO'],order=(0,1,0),seasonal_order=(
results = model.fit()
print(results.summary())
df_CSCO['Forecast'] = results.predict()
df_CSCO[['NASDAQ.CSCO','Forecast']].plot(figsize=(20,8))
```

C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:225: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.


' ignored when e.g. forecasting.', ValueWarning)

Statespace Model Results

```
=====
=====
Dep. Variable:          NASDAQ.CSCO    No. Observations:
41265
Model:                SARIMAX(0, 1, 0)x(0, 1, 0, 12)    Log Likelihood          8
5502.595
Date:                Mon, 25 Feb 2019    AIC          -17
1003.190
Time:                15:05:33    BIC          -17
0994.563
Sample:                0    HQIC          -17
1000.463
                        - 41265
Covariance Type:      opg
=====
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
sigma2          0.0009    1.54e-07    6012.819      0.000      0.001      0.001
=====
=
Ljung-Box (Q):          11736.64    Jarque-Bera (JB):          21073382447.0
0
Prob(Q):                0.00    Prob(JB):                0.0
0
Heteroskedasticity (H): 0.30    Skew:                2.6
7
Prob(H) (two-sided):    0.00    Kurtosis:          3504.4
6
=====
=
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



In [94]:

Out[94]:

	NASDAQ.CSCO	First_Difference	Forecast
Month			
1970-01-01	33.8800	0.1400	0.0000
1970-01-01	33.9000	0.0200	33.8800
1970-01-01	33.8499	-0.0501	33.9000
1970-01-01	33.8400	-0.0099	33.8499
1970-01-01	33.8800	0.0400	33.8400

In [95]:

C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.
ValueWarning)

Out[95]: 41265 32.225
41266 32.190
41267 32.170
41268 32.150
41269 32.180
41270 32.170
41271 32.150
41272 32.165
41273 32.180
41274 32.180
dtype: float64

In [96]:

C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.
ValueWarning)

Out[96]: 41264 32.195
41265 32.225
41266 32.190
41267 32.170
41268 32.150
41269 32.180
41270 32.170
41271 32.150
41272 32.165
41273 32.180
41274 32.180
41275 32.175
dtype: float64

In [97]:

```
from sklearn.metrics import mean_squared_error, mean_absolute_error
print('Mean Squared Error NASDAQ.CSCO -', mean_squared_error(df_CSCO['NASDAQ.CSCO'], df_CSCO['NASDAQ.CSCO']))



Mean Squared Error NASDAQ.CSCO - 0.0356937844969608
Mean Absolute Error NASDAQ.CSCO - 0.015775407730929027
```

```
In [98]: #Time Series Forecasting for NASDAQ.EBAY

df_EBAY = final[['Month',stock_features[4]]]
print(df_EBAY.head())
df_EBAY.set_index('Month',inplace=True)
print(df_EBAY.head())
df_EBAY.plot()
plt.title("Time Series Plot for NASDAQ.EBAY")
plt.show()
#Test Staionarity
```

	Month	NASDAQ.EBAY
0	1970-01-01	33.3975
1	1970-01-01	33.3950
2	1970-01-01	33.4100
3	1970-01-01	33.3350
4	1970-01-01	33.4000


	Month	NASDAQ.EBAY
1970-01-01	33.3975	
1970-01-01	33.3950	
1970-01-01	33.4100	
1970-01-01	33.3350	
1970-01-01	33.4000	



```
Augmented Dickey-Fuller Test:
ADF Test Statistic : -1.8757616359415816
p-value : 0.343548087802396
#Lags Used : 47
Number of Observations Used : 41218
Critical 1% : value -3.430508661441506
```

```
In [99]: #MAKING TIME SERIES STATIONARY
#Differencing


df_EBAY = df_EBAY.copy()
df_EBAY['First_Difference'] = df_EBAY['NASDAQ.EBAY'] - df_EBAY['NASDAQ.EBAY'].shift(1)
df_EBAY.dropna(inplace=True)
#test Stationarity
```




```
Augmented Dickey-Fuller Test:
ADF Test Statistic : -1.8639133106593535
p-value : 0.34922311499828906
#Lags Used : 47
Number of Observations Used : 41217
Critical 1% : value -3.4305086652911636
Critical 5% : value -2.8616101264708296
Critical 10% : value -2.5668073260584587
weak evidence against null hypothesis, time series has a unit root, indicating it
is non-stationary
```

```
In [100]: #Seasonal Decomposition
from statsmodels.tsa.seasonal import seasonal_decompose
plt.figure(figsize=(11,8))
decomposition = seasonal_decompose(df_EBAY['NASDAQ.EBAY'],freq=12)
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid
plt.subplot(411)
plt.plot(df_EBAY['NASDAQ.EBAY'],label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend,label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal,label='Seasonal')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual,label='Residual')
```

```
Out[100]: <matplotlib.legend.Legend at 0x1b6a3819588>
```



```
In [101]: ts_log_decompose = residual
ts_log_decompose.dropna(inplace=True)
```



Augmented Dickey-Fuller Test:

ADF Test Statistic : -44.88049175892131

p-value : 0.0

#Lags Used : 55

Number of Observations Used : 41197

Critical 1% : value -3.4305087423235587

Critical 5% : value -2.861610160516496

Critical 10% : value -2.566807344180027


strong evidence against the null hypothesis, reject the null hypothesis. Data has no unit root and is stationary

```
In [102]: # Note : This is stationary because:

#Test statistic is lower than critical values.
```

```
In [103]: #Autocorealtion plot and Partial Autocorelation plots

fig = plt.figure(figsize=(20,8))
ax1 = fig.add_subplot(211)
fig = sm.graphics.tsa.plot_acf(df_EBAY['First_Difference'].iloc[47:], lags=80, ax=ax1)
ax2 = fig.add_subplot(212)
```



```
In [104]: lag_acf = acf(df_EBAY['First_Difference'],nlags=80)
```

```

In [105]: plt.figure(figsize=(20,8))
plt.subplot(121)
plt.plot(lag_acf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_EBAY['First_Difference'])),linestyle='--',color='g')
plt.axhline(y=1.96/np.sqrt(len(df_EBAY['First_Difference'])),linestyle='--',color='g')

plt.title('Autocorrelation')


plt.subplot(122)

plt.plot(lag_pacf)
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df_EBAY['First_Difference'])),linestyle='--',color='g')
plt.axhline(y=1.96/np.sqrt(len(df_EBAY['First_Difference'])),linestyle='--',color='g')

plt.title('Partial Autocorrelation')

plt.tight_layout()

```



```

In [106]: #Note- The two dotted lines on either sides of 0 are the confidence intervals.

#These can be used to determine the 'p' and 'q' values as:

#p: The first time where the PACF crosses the upper confidence interval, here its clo

```

```
In [107]: # fit model
model= sm.tsa.statespace.SARIMAX(df_EBAY['NASDAQ.EBAY'],order=(0,1,0),seasonal_order=
results = model.fit()
print(results.summary())
df_EBAY['Forecast'] = results.predict()
df_EBAY[['NASDAQ.EBAY', 'Forecast']].plot(figsize=(20,8))
```

C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:225: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.


' ignored when e.g. forecasting.', ValueWarning)

Statespace Model Results

```
=====
=====
Dep. Variable:                NASDAQ.EBAY    No. Observations:
41265
Model:                SARIMAX(0, 1, 0)x(0, 1, 0, 12)    Log Likelihood            8
2104.712
Date:                Mon, 25 Feb 2019    AIC            -16
4207.424
Time:                15:08:07    BIC            -16
4198.797
Sample:                0    HQIC            -16
4204.697
                        - 41265
Covariance Type:                opg
=====
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
sigma2          0.0011    9.43e-07   1158.859    0.000        0.001        0.001
=====
=
Ljung-Box (Q):                10939.63    Jarque-Bera (JB):                28223015.4
4
Prob(Q):                0.00    Prob(JB):                0.0
0
Heteroskedasticity (H):                1.21    Skew:                0.3
5
Prob(H) (two-sided):                0.00    Kurtosis:                131.1
4
=====
=====
=
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



In [108]:

Out[108]:

	NASDAQ.EBAY	First_Difference	Forecast
Month			
1970-01-01	33.395	-0.0025	0.000
1970-01-01	33.410	0.0150	33.395
1970-01-01	33.335	-0.0750	33.410
1970-01-01	33.400	0.0650	33.335
1970-01-01	33.430	0.0300	33.400

```
In [109]: from sklearn.metrics import mean_squared_error, mean_absolute_error
print('Mean Squared Error NASDAQ.EBAY -', mean_squared_error(df_EBAY['NASDAQ.EBAY'], d
```

```
Mean Squared Error NASDAQ.EBAY - 0.03483567893982985
Mean Absolute Error NASDAQ.EBAY - 0.02168803344281537
```

In [110]:

```
C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: Val
ueWarning: No supported index is available. Prediction results will be given with
an integer index beginning at `start`.
ValueWarning)
```

```
Out[110]: 41265    36.090
41266    36.030
41267    36.030
41268    36.020
41269    36.020
41270    36.025
41271    36.020
41272    36.025
41273    36.020
41274    36.020
dtype: float64
```

In [111]:

```
C:\Users\HP\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:531: Val
ueWarning: No supported index is available. Prediction results will be given with
an integer index beginning at `start`.
ValueWarning)
```

```
Out[111]: 41265    36.090
41266    36.030
41267    36.030
41268    36.020
41269    36.020
41270    36.025
41271    36.020
41272    36.025
41273    36.020
41274    36.020
41275    36.010
dtype: float64
```

CONCLUSION : The predicted stock prices values have been stored in the Forecast Columns of the each stock entity dataframe

